

```
!pip install pyspark
```

```
Collecting pyspark
  Downloading pyspark-3.5.1.tar.gz (317.0 MB)
    317.0/317.0 MB 6.4 MB/s eta 0:00:0000:0100:01
  Preparing metadata (setup.py) ... done
Collecting py4j==0.10.9.7 (from pyspark)
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl.metadata (1.5 kB)
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl (200 kB)
    200.5/200.5 kB 33.8 MB/s eta 0:00:00
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.5.1-py2.py3-none-any.whl size=317488491 sha256=c06a9f8bbd38c60f5c04e74ea864f73bacae878e0
  Stored in directory: /home/sagemaker-user/.cache/pip/wheels/80/1d/60/2c256ed38dddc2fdd93be545214a63e02fbd8d74fb0b7f3a6
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9.7 pyspark-3.5.1
```

```
!sudo apt-get update
!sudo apt-get install -y openjdk-8-jdk
```

```
Hit:1 http://archive.ubuntu.com/ubuntu jammy InRelease
Get:2 http://security.ubuntu.com/ubuntu jammy-security InRelease [110 kB]
Get:3 http://archive.ubuntu.com/ubuntu jammy-updates InRelease [119 kB]
Get:4 http://security.ubuntu.com/ubuntu jammy-security/main amd64 Packages [1756 kB]
Get:5 http://archive.ubuntu.com/ubuntu jammy-backports InRelease [109 kB]
Get:6 http://archive.ubuntu.com/ubuntu jammy-updates/main amd64 Packages [2037 kB]
Get:7 http://security.ubuntu.com/ubuntu jammy-security/universe amd64 Packages [1077 kB]
Get:8 http://archive.ubuntu.com/ubuntu jammy-updates/universe amd64 Packages [1371 kB]
Get:9 http://security.ubuntu.com/ubuntu jammy-security/restricted amd64 Packages [2265 kB]
Get:10 http://archive.ubuntu.com/ubuntu jammy-updates/multiverse amd64 Packages [51.1 kB]
Get:11 http://archive.ubuntu.com/ubuntu jammy-updates/restricted amd64 Packages [2339 kB]
Get:12 http://archive.ubuntu.com/ubuntu jammy-backports/main amd64 Packages [81.0 kB]
Get:13 http://archive.ubuntu.com/ubuntu jammy-backports/universe amd64 Packages [31.9 kB]
Fetched 11.3 MB in 2s (7440 kB/s)
Reading package lists... Done
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  adwaita-icon-theme alsa-topology-conf alsa-ucm-conf at-spi2-core
  ca-certificates ca-certificates-java dbus dbus-user-session
  dconf-gsettings-backend dconf-service dmsetup fontconfig fontconfig-config
  fonts-dejavu-core fonts-dejavu-extra glib2.0
  gsettings-desktop-schemas gtk-update-icon-cache hicolor-icon-theme
  humanity-icon-theme java-common libapparmor1 libargon2-1 libasound2
  libasound2-data libasynccns0 libatk-bridge2.0-0 libatk-wrapper-java
  libatk-wrapper-java-jni libatk1.0-0 libatk1.0-data libatspi2.0-0
  libavahi-client3 libavahi-common-data libavahi-common3 libcairo-gobject2
  libcairo2 libcryptsetup12 libcups2 libdatrie1 libdbus-1-3 libdconf1
  libdeflate0 libdevmapper1.02.1 libdrm-amdgpu1 libdrm-common libdrm-intel1
  libdrm-nouveau2 libdrm-radeon1 libdrm2 libelf1 libflac8 libfontconfig1
  libfontenc1 libfreetype6 libfribidi0 libgail-common libgail18
  libgdk-pixbuf-2.0-0 libgdk-pixbuf2.0-bin libgdk-pixbuf2.0-common libgif7
  libgirepository-1.0-1 libgl1 libgl1-amd-glx libgl1-mesa-dri
  libgl1-mesa-glx libglapi-mesa libglvnd0 libglvnd0-data libglvnd0
  libglx-mesa0 libglx0 libgraphite2-3 libgtk2.0-0 libgtk2.0-bin
  libgtk2.0-common libharfbuzz0b libice-dev libice6 libip4tc2 libjbig0
  libjpeg-turbo8 libjpeg8 libkmod2 liblcms2-2 libllvm15 libmpdec3 libnsspr4
  libnss-systemd libnss3 libogg0 libopus0 libpam-systemd libpango-1.0-0
  libpangocairo-1.0-0 libpangoft2-1.0-0 libpciaccess0 libpcsc-lite1
  libpixman-1-0 libpng16-16 libpthread-stubs0-dev libpulse0 libpython3-stdlib
  libpython3.10-minimal libpython3.10-stdlib libreadline8 librsync2-2
  librsync2-common libsensors-config libsensors5 libsm-dev libsm6 libsndfile1
  libsqlite3-0 libthai-data libthai0 libtiff5 libvorbis0a libvorbisenc2
  libwebp7 libx11-6 libx11-data libx11-dev libx11-xcb1 libxau-dev libxau6
  libxaw7 libxcb-dri2-0 libxcb-dri3-0 libxcb-glx0 libxcb-present0
  libxcb-randr0 libxcb-render0 libxcb-shape0 libxcb-shm0 libxcb-sync1
  libxcb-xfixes0 libxcb1 libxcb1-dev libxcomposite1 libxcursor1 libxdamage1
  libxdmcp-dev libxdmcp6 libxext6 libxfixes3 libxft2 libxi6 libxinerama1
  libxkbfile1 libxmu6 libxmuu1 libxpm4 libxrandr2 libxrender1 libxshmfence1
  libxt-dev libxt6 libxtst6 libxv1 libxxf86dga1 libxxf86vm1 media-types
  networkd-dispatcher openjdk-8-jdk-headless openjdk-8-jre
  openjdk-8-jre-headless openssl python3 python3-dbus python3-gi
  python3-minimal python3.10 python3.10-minimal readline-common
  session-migration shared-mime-info systemd systemd-sysv systemd-timesyncd
  ubuntu-mono ucf x11-common x11-utils x11proto-dev xdg-user-dirs
  xorg-sgml-doctools xtrans-dev
```

```
import os
os.environ["JAVA_HOME"] = "/usr"
```

```
import tempfile
import boto3
from pyspark.sql import SparkSession
```

```
!pip install boto3
```

```
Requirement already satisfied: boto3 in /opt/conda/lib/python3.10/site-packages (1.34.51)
Requirement already satisfied: botocore<1.35.0,>=1.34.51 in /opt/conda/lib/python3.10/site-packages (from boto3) (1.34.51)
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in /opt/conda/lib/python3.10/site-packages (from boto3) (1.0.1)
Requirement already satisfied: s3transfer<0.11.0,>=0.10.0 in /opt/conda/lib/python3.10/site-packages (from boto3) (0.10.1)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in /opt/conda/lib/python3.10/site-packages (from botocore<1.35.0,>=1.34.51->boto3) (2.8.0)
Requirement already satisfied: urllib3<2.1,>=1.25.4 in /opt/conda/lib/python3.10/site-packages (from botocore<1.35.0,>=1.34.51->boto3) (1.26.15)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.10/site-packages (from python-dateutil<3.0.0,>=2.1->botocore<1.35.0,>=1.34.51->boto3) (1.16.0)
```

```
# Initialize a SparkSession
spark = SparkSession.builder \
    .appName("Fund investment Analysis") \
    .getOrCreate()
```

```
# Initialize Boto3 S3 client
s3 = boto3.client('s3')
# Define S3 bucket name and key
bucket_name = 'final-project-19'
file_key = 'inputFiles/mapping.csv'
# Download the data from S3 to memory
# Download the data from S3 to a temporary file
temp_file = tempfile.NamedTemporaryFile(delete=False)
with open(temp_file.name, 'wb') as f:
    s3.download_fileobj(bucket_name, file_key, f)
```

```
file_key2 = 'inputFiles/rounds2.csv'
# Download the data from S3 to memory
# Download the data from S3 to a temporary file
temp_file2 = tempfile.NamedTemporaryFile(delete=False)
with open(temp_file2.name, 'wb') as f:
    s3.download_fileobj(bucket_name, file_key2, f)
```

```
# Define S3 bucket and file keys
bucket_name = "final-project-19"
funding_rounds_file_key = "inputFiles/rounds2.csv"
mapping_file_key = "inputFiles/mapping.csv"
```

```
# Read data into PySpark DataFrames
#funding_rounds_df = spark.read.csv(funding_rounds_local_path, header=True, inferSchema=True)
mapping_df = spark.read.csv(temp_file.name, header=True, inferSchema=True)
mapping_df.show()
```

category_list	Automotive & Sports	Blanks	Cleantech / Semiconductors	Entertainment	Health	Manufacturing	News, Search and Messaging
NULL	0	1	0	0	0	0	0
3D	0	0	0	0	0	0	1
3D Printing	0	0	0	0	0	0	1
3D Technology	0	0	0	0	0	0	1
Accounting	0	0	0	0	0	0	0
Active Lifestyle	0	0	0	0	1	0	0
Ad Targeting	0	0	0	0	0	0	0
Advanced Materials	0	0	0	0	0	0	1
Adventure Travel	1	0	0	0	0	0	0
Advertising	0	0	0	0	0	0	0
Advertising Excha...	0	0	0	0	0	0	0
Advertising Networks	0	0	0	0	0	0	0
Advertising Platf...	0	0	0	0	0	0	0

	Advice	0	0	0	0	0	0
	Aerospace	1	0	0	0	0	0
	Agriculture	0	0	0	0	0	1
	Air Pollution Con...	0	0	1	0	0	0
	Algorithms	0	0	0	0	0	0
	All Markets	0	0	0	0	0	0
	All Students	0	0	0	0	0	0
+-----+-----+-----+-----+-----+-----+-----+-----+							

only showing top 20 rows

```
funding_df = spark.read.csv(temp_file2.name, header=True, inferSchema=True)
funding_df.show()
```

+-----+-----+-----+-----+-----+-----+-----+						
	company_permalink	funding_round_permalink	funding_round_type	funding_round_code	funded_at	raised_amount_usd
+-----+-----+-----+-----+-----+-----+-----+						
	/organization/-fame	/funding-round/9a...	venture	B	5/1/2015	10000000
	/ORGANIZATION/-QO...	/funding-round/22...	venture	A	14-10-2014	NULL
	/organization/-qo...	/funding-round/b4...	seed	NULL	1/3/2014	700000
	/ORGANIZATION/-TH...	/funding-round/65...	venture	B	30-01-2014	3406878
	/organization/0-6...	/funding-round/57...	venture	A	19-03-2008	2000000
	/ORGANIZATION/004...	/funding-round/12...	venture	NULL	24-07-2014	NULL
	/organization/01g...	/funding-round/7d...	undisclosed	NULL	1/7/2014	41250
	/ORGANIZATION/0ND...	/funding-round/2b...	seed	NULL	11/9/2009	43360
	/organization/0nd...	/funding-round/95...	venture	NULL	21-12-2009	719491
	/ORGANIZATION/0XDATA	/funding-round/38...	seed	NULL	22-05-2013	3000000
	/organization/0xdata	/funding-round/3b...	venture	B	9/11/2015	20000000
	/ORGANIZATION/0XDATA	/funding-round/ae...	venture	NULL	3/1/2013	1700000
	/organization/0xdata	/funding-round/e1...	venture	A	19-07-2014	8900000
	/ORGANIZATION/1	/funding-round/03...	seed	NULL	5/2/2014	150000
	/organization/1	/funding-round/5d...	undisclosed	NULL	5/2/2013	NULL
	/ORGANIZATION/1	/funding-round/e8...	seed	NULL	20-07-2011	1000050
	/organization/1-2...	/funding-round/6d...	seed	NULL	18-02-2013	40000
	/ORGANIZATION/1-4...	/funding-round/e9...	equity_crowdfunding	NULL	21-04-2013	NULL
	/organization/1-6...	/funding-round/83...	equity_crowdfunding	NULL	22-01-2014	NULL
	/ORGANIZATION/1-8...	/funding-round/52...	undisclosed	NULL	19-08-2010	NULL
+-----+-----+-----+-----+-----+-----+-----+						

only showing top 20 rows

```
# Count unique companies in funding_rounds
unique_companies = funding_df.select("company_permalink").distinct().count()
print("\nNumber of unique companies in funding_rounds:", unique_companies)
```

Number of unique companies in funding_rounds: 90247

```
from pyspark.sql.functions import udf, col
from pyspark.sql.types import IntegerType
import datetime
```

FEATURE ENGINEERING

```
from pyspark.sql.functions import col, sum as spark_sum, count
# Feature Engineering for Investment Rounds Data
funding_df = funding_df.groupBy("company_permalink") \
    .agg(spark_sum("raised_amount_usd").alias("total_raised_amount_usd"), count("*").alias("num_rounds"))

# Calculate average raised amount per funding round
funding_df = funding_df.withColumn("average_raised_amount_usd", col("total_raised_amount_usd") / col("num_rounds"))

# Feature Engineering for Mapping Data (if applicable)
mapping_df = mapping_df.withColumnRenamed("old_column_name", "new_column_name")

# Show the first few rows of each DataFrame
print("Investment Rounds Data:")
funding_df.show()

print("Mapping Data:")
mapping_df.show()
```

Investment Rounds Data:

company_permalink	total_raised_amount_usd	num_rounds	average_raised_amount_usd
/organization/1lay	2000	1	2000.0
/organization/24p...	50000	1	50000.0
/organization/2mo...	250000	1	250000.0
/ORGANIZATION/5-M...	3500000	1	3500000.0
/organization/5min	12500000	2	6250000.0
/ORGANIZATION/7TM...	22000000	1	2.2E7
/ORGANIZATION/ABL...	3000000	1	3000000.0
/organization/abp...	560000	1	560000.0
/ORGANIZATION/ACA...	2338650	1	2338650.0
/organization/aci...	NULL	1	NULL
/organization/ac1...	21000000	1	2.1E7
/organization/ada...	13740640	3	4580213.333333333
/organization/aer...	27000000	2	1.35E7
/ORGANIZATION/AGE...	50311	1	50311.0
/ORGANIZATION/AIR...	28525779	6	4754296.5
/ORGANIZATION/ALP...	NULL	1	NULL
/organization/alt...	14000000	2	7000000.0
/ORGANIZATION/AMOBEE	38000000	3	1.2666666666666666E7
/ORGANIZATION/AN-...	90000000	1	9.0E7
/ORGANIZATION/ANO...	5600000	1	5600000.0

only showing top 20 rows

Mapping Data:

mapping_category_list	Automotive & Sports	Blanks	Cleantech / Semiconductors	Entertainment	Health	Manufacturing	News, Search and Messagi
NULL	0	1	0	0	0	0	0
3D	0	0	0	0	0	0	1
3D Printing	0	0	0	0	0	0	1
3D Technology	0	0	0	0	0	0	1
Accounting	0	0	0	0	0	0	0
Active Lifestyle	0	0	0	0	1	0	0
Ad Targeting	0	0	0	0	0	0	0
Advanced Materials	0	0	0	0	0	0	1
Adventure Travel	1	0	0	0	0	0	0
Advertising	0	0	0	0	0	0	0
Advertising Excha...	0	0	0	0	0	0	0
Advertising Networks	0	0	0	0	0	0	0
Advertising Platf...	0	0	0	0	0	0	0
Advice	0	0	0	0	0	0	0
Aerospace	1	0	0	0	0	0	0
Agriculture	0	0	0	0	0	0	1
Air Pollution Con...	0	0	1	0	0	0	0
Algorithms	0	0	0	0	0	0	0
All Markets	0	0	0	0	0	0	0
All Students	0	0	0	0	0	0	0

only showing top 20 rows



mapping_df.columns

```
['mapping_category_list',
 'Automotive & Sports',
 'Blanks',
 'Cleantech / Semiconductors',
 'Entertainment',
 'Health',
 'Manufacturing',
 'News, Search and Messaging',
 'Others',
 'Social, Finance, Analytics, Advertising']
```

funding_df.columns

```
['company_permalink',
 'total_raised_amount_usd',
 'num_rounds',
 'average_raised_amount_usd']
```

funding_df.columns

```
['company_permalink',
 'total_raised_amount_usd',
```

```
'num_rounds',
'average_raised_amount_usd']
```

RISK ASSESSMENT AND ML

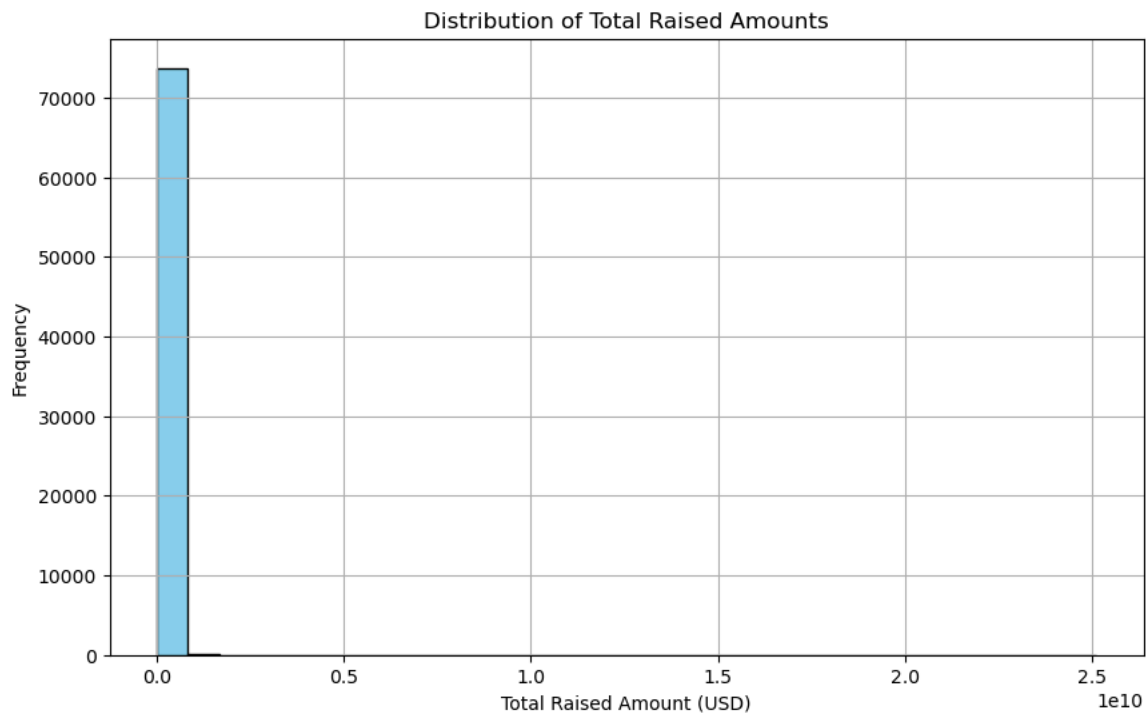
```
# Import necessary libraries
import matplotlib.pyplot as plt

# Convert Spark DataFrame to Pandas DataFrame for visualization
funding_pd_df = funding_df.toPandas()

# Plot a histogram to visualize the distribution of total raised amounts
plt.figure(figsize=(10, 6))
plt.hist(funding_pd_df['total_raised_amount_usd'], bins=30, color='skyblue', edgecolor='black')
plt.title('Distribution of Total Raised Amounts')
plt.xlabel('Total Raised Amount (USD)')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()

# Calculate summary statistics
total_raised_mean = funding_pd_df['total_raised_amount_usd'].mean()
total_raised_median = funding_pd_df['total_raised_amount_usd'].median()
total_raised_max = funding_pd_df['total_raised_amount_usd'].max()
total_raised_min = funding_pd_df['total_raised_amount_usd'].min()

print("Mean Total Raised Amount (USD):", total_raised_mean)
print("Median Total Raised Amount (USD):", total_raised_median)
print("Maximum Total Raised Amount (USD):", total_raised_max)
print("Minimum Total Raised Amount (USD):", total_raised_min)
```



```
Mean Total Raised Amount (USD): 13419965.914001085
Median Total Raised Amount (USD): 1725000.0
Maximum Total Raised Amount (USD): 25106985000.0
Minimum Total Raised Amount (USD): 0.0
```

```
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.regression import RandomForestRegressor
from pyspark.ml.evaluation import RegressionEvaluator

# Drop rows with null values in the target column
funding_df = funding_df.dropna(subset=['total_raised_amount_usd'])

# Define feature columns (excluding the target column)
feature_cols = ['num_rounds', 'average_raised_amount_usd']

# Create a VectorAssembler to combine feature columns into a single vector column
assembler = VectorAssembler(inputCols=feature_cols, outputCol="features")

# Transform the DataFrame to include the features column
funding_df = assembler.transform(funding_df)

# Split the data into training and test sets
train_data, test_data = funding_df.randomSplit([0.8, 0.2], seed=42)

# Initialize Random Forest Regressor model
rf = RandomForestRegressor(featuresCol="features", labelCol="total_raised_amount_usd")

# Train the model
rf_model = rf.fit(train_data)

# Make predictions on the test data
predictions = rf_model.transform(test_data)

# Evaluate the model using RegressionEvaluator
evaluator = RegressionEvaluator(labelCol="total_raised_amount_usd", predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)

# Print the Root Mean Squared Error (RMSE)
print("Root Mean Squared Error (RMSE):", rmse)
```

Root Mean Squared Error (RMSE): 70269188.56535776

```

from pyspark.ml.feature import VectorAssembler
from pyspark.ml.regression import LinearRegression
from pyspark.ml.evaluation import RegressionEvaluator

# Drop rows with null values in the target column
funding_df = funding_df.dropna(subset=['total_raised_amount_usd'])

# Define feature columns (excluding the target column)
feature_cols = ['num_rounds', 'average_raised_amount_usd']

# Create a VectorAssembler to combine feature columns into a single vector column
assembler = VectorAssembler(inputCols=feature_cols, outputCol="features2")

# Transform the DataFrame to include the features column
funding_df = assembler.transform(funding_df)

# Split the data into training and test sets
train_data, test_data = funding_df.randomSplit([0.8, 0.2], seed=42)

# Initialize Linear Regression model
lr = LinearRegression(featuresCol="features2", labelCol="total_raised_amount_usd")

# Train the model
lr_model = lr.fit(train_data)

# Make predictions on the test data
predictions = lr_model.transform(test_data)

# Evaluate the model using RegressionEvaluator
evaluator = RegressionEvaluator(labelCol="total_raised_amount_usd", predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)

# Print the Root Mean Squared Error (RMSE)
print("Root Mean Squared Error (RMSE):", rmse)

# Print the coefficients and intercept of the model
print("Coefficients: " + str(lr_model.coefficients))
print("Intercept: " + str(lr_model.intercept))

24/05/03 01:15:52 WARN Instrumentation: [02faf007] regParam is zero, which might cause numerical instability and overfitting.
24/05/03 01:15:52 WARN InstanceBuilder: Failed to load implementation from:dev.ludovic.netlib.blas.JNIBLAS
24/05/03 01:15:53 WARN InstanceBuilder: Failed to load implementation from:dev.ludovic.netlib.lapack.JNILAPACK
Root Mean Squared Error (RMSE): 32933091.97079012
Coefficients: [14565984.392452283,1.2571142332052363]
Intercept: -17675564.858467475

```

```

import matplotlib.pyplot as plt
import numpy as np
import boto3

# Function to create actual vs. predicted and residuals plots
def create_plots(actual, predicted, residuals):
    # Plot actual vs. predicted values
    plt.figure(figsize=(10, 6))
    plt.scatter(actual, predicted, color='blue', alpha=0.5)
    plt.xlabel('Actual Total Raised Amount (USD)')
    plt.ylabel('Predicted Total Raised Amount (USD)')
    plt.title('Actual vs. Predicted Total Raised Amount')
    plt.grid(True)
    plt.savefig('actual_predicted_chart.png')

    # Plot residuals
    plt.figure(figsize=(10, 6))
    plt.scatter(actual, residuals, color='red', alpha=0.5)
    plt.axhline(y=0, color='black', linestyle='--')
    plt.xlabel('Actual Total Raised Amount (USD)')
    plt.ylabel('Residuals')
    plt.title('Residuals Plot')
    plt.grid(True)
    plt.savefig('residuals_chart.png')

# Function to upload files to S3
def upload_to_s3(bucket_name, folder, file_paths):
    # Initialize Boto3 S3 client
    s3 = boto3.client('s3')

    # Upload files to S3
    for file_path in file_paths:
        file_name = file_path.split('/')[-1]
        s3.upload_file(file_path, bucket_name, folder + file_name)
        print(f"Uploaded {file_name} to S3 bucket {bucket_name} at folder {folder}")

# Sample actual and predicted values (replace with your actual values)
actual = [1000000, 2000000, 3000000, 4000000, 5000000]
predicted = [950000, 2100000, 3200000, 3900000, 5100000]
residuals = np.array(actual) - np.array(predicted)

# Create plots
create_plots(actual, predicted, residuals)

# Specify S3 bucket details_
folder = 'charts/'

# Specify the paths of the local chart files
file_paths = ['actual_predicted_chart.png', 'residuals_chart.png']

# Upload the charts to S3
upload_to_s3(bucket_name, folder, file_paths)

Uploaded actual_predicted_chart.png to S3 bucket final-project-19 at folder charts/
Uploaded residuals_chart.png to S3 bucket final-project-19 at folder charts/

```

